

Name	Description	Default
rhipe_map_buffsize	The length of map.keys and map.values	10,000
rhipe_reduce_buffsize	The length of reduce.values	10,000
io.sort.mb	The number of megabytes of intermediate output before it is locally reduced when combiner is true	Hadoop site file
mapred.textoutputformat.separator	Character that separates the key and value in text output	[tab]
mapred.field.separator	Character that separates the elements of vectors in text output	[space]
mapred.textoutputformat.usekey	For text output, if false, the key is not written	TRUE
mapred.map.tasks	The number of map tasks to divide the job into	For sequence and text file input, it is derived from the number of blocks. For lapply input, it is derived from N in the call to <code>rhmr</code>
mapred.reduce.tasks	The number of reduce tasks to launch to aggregate the intermediate keys	Hadoop site file
mapred.tasktimeout	The number of minutes after which the Tasktracker will terminate the R process. Set to zero, if the job consists of unpredictable long running tasks	10 min
rhipe_stream_buffer	The size of the buffer in KB for transferring data between Java and R	1MB
mapred.input.filename	the name of current input file for the given task, read it using <code>Sys.getenv</code>	